

Latency distribution: cache_hot vs cache_cold (5 runs)

